

PTS — Two-Score Framework

An evidence-based tool that assesses a single publication on two independent 100-point scores: **PTS-A** (Antisemitism, IHRA) and **PTS-J** (Journalistic Standards, IMPRESS). A model returns per-criterion judgements with exact quotes; a deterministic scorer verifies every quote and applies the rules. The two scores are never blended.

PTS-A · IHRA · AS1-AS6

PTS-J · IMPRESS · J1-J11 + JD

Deterministic scorer

Quote verification

Stance gate

Human-review routing

How it works

Give the app an article by URL or pasted text. It extracts the publication into labelled sections (`HEADLINE` , `STANDFIRST` , `BYLINE` , `PUBLISHED` , `SOURCE` , `BODY`), sends it to a model with a **forced, structured tool call**, then runs a **deterministic scorer** that verifies every piece of evidence and computes both numbers. The model never computes the score — software does.

Providers, in order: **OpenRouter** (if `OPENROUTER_API_KEY` is set; default `anthropic/claude-haiku-4.5`) → **Anthropic** (if `ANTHROPIC_API_KEY` is set; default `claude-haiku-4.5`) → **Mock** (labelled all-PASS, no key needed).

PTS-A — Antisemitism Score (IHRA), 100 points

ID	POINTS	WHAT IT CHECKS	IHRA
AS1	20 crit	Violence & incitement — calling for, aiding, justifying or celebrating harm to Jews, Jewish institutions or property.	1
AS2	20	Collective tropes & imagery — Jewish power/control/conspiracy over media, finance, government; dehumanising stereotypes; classic imagery.	2
AS3	15	Collective responsibility & dual loyalty — blaming Jews as a people for acts of individuals or of Israel; dual-loyalty accusations.	3, 6, 11
AS4	20 crit	Holocaust denial & distortion — denying fact/scope/mechanisms/intent; claiming it was invented or exaggerated.	4, 5
AS5	15 cont	Israel-related — within legitimate criticism: no delegitimation ("racist endeavour"), Nazi comparisons, blood-libel imagery, or double standards.	7, 8, 9, 10
AS6	10	Amplification & gratuitous identity — third-party antisemitic content unattributed/unchallenged; irrelevant "Jewish" labelling.	EJN; 4.2

crit = own-voice failure caps PTS-A at 39 & forces review · **cont** = contested zone, always human review.

PTS-J — Journalistic Standards Score (IMPRESS), 100 points

IMPRESS Clause 4 (Discrimination) is **not** in PTS-J — it lives entirely in PTS-A. 80 criteria points (J1-J11) + a 20-point conduct block (JD).

ID	POINTS	WHAT IT CHECKS	CLAUSE
J1	20	Accuracy — no significant inaccuracy, misrepresentation or distortion (<code>fabrication</code> flag for invented facts/quotes).	1.1 / 1.4
J2	10	Fact clearly distinguished from opinion and conjecture.	1.3
J3	8	Significant errors corrected with due prominence.	1.2

ID	POINTS	WHAT IT CHECKS	CLAUSE
J4	7	Headline / standfirst / preview do not misrepresent the body.	1.5 *
J5	7	Third-party material attributed; no plagiarism.	2.1
J6	3	Failure to credit, corrected.	2.2
J7	8	Conflicts of interest & ownership disclosed (EMFA Art. 6).	10.2
J8	6	Paid or sponsored content clearly labelled.	10.1
J9	3	Financial information objective, interests disclosed.	10.3
J10	3	Failure to disclose, corrected.	10.4
J11	5	Public-interest justification evident where relied upon.	—
JD	20	Conduct — clauses 3 (Children), 5 (Harassment), 6 (Justice), 7 (Privacy), 8 (Sources), 9 (Suicide). Starts at 20; –10 per verified engaged-and-breached clause (floor 0).	3,5,6,7,8,9

* J4 is a PTS extension of clause 1.5.

Cross-cutting scoring rules

Evidence verification

Every **FAIL** needs an exact quote, checked against the text with tolerant matching (whitespace / curly-quote normalisation, ellipsis splitting, label stripping, approximate match ≥ 0.85). No verifiable quote → **no deduction**; the finding is rejected and queued for review.

Stance gate (PTS-A only)

AS1-AS5 can fail only for **OWN_VOICE** or **UNCritical_AMPLIFICATION**. **REPORTED_CONTEXTUALISED** / **COUNTERED** content is not penalised. A **FAIL** with **failure_stance: NONE** inherits **overall_stance**.

Critical caps (39)

AS1 or **AS4** failing in the publication's **own voice** caps PTS-A at 39. **Fabrication** (**J1** , or a fabricated source under conduct clause 8) caps PTS-J at 39.

NOT_ASSESSABLE & rescaling

Unassessable criteria leave the "possible" total and the score is rescaled to 100 — so coverage can be below 100 while the score is 100. The **absence** of a problem is a **PASS**, not **NOT_ASSESSABLE**.

Human-review routing

AS1 , **AS4** , **AS5** failures; any model-flagged review; and any accepted failure with confidence < 0.5 are routed to a de-duplicated review queue (per score).

Legal flag

A separate routing flag (**INCITEMENT** / **HOLOCAUST_DENIAL**) — never a legal finding, and never part of either score.

Designation profiles & tier bands

Designations

- **ARTICLE / DOCUMENTARY** — all AS + all J + JD.
- **POST** — PTS-A in full; PTS-J only on **J1** , **J2** , **J5** + JD.
- **SATIRE** — labelled, **not scored** (both scores **null**).

Tier bands

PTS-A: ≥ 100 No indicators · ≥ 75 Minor concerns · ≥ 60 Amplification/contested · ≥ 40 Antisemitic content present · else Severe.

PTS-J: ≥ 100 Compliant · ≥ 75 Generally compliant · ≥ 60 Compliant with exceptions · ≥ 40 Breach · else Serious breach.

Scope & limitations

These scores judge the text as written. They detect antisemitism (IHRA) and journalistic-standards breaches (IMPRESS) that are visible in the piece itself, including how claims are attributed.

They are **not** a fact-check of claims against external reality, and **not** a bias, balance or one-sidedness meter. Properly attributed, reported content (a **REPORTED_CONTEXTUALISED** stance) is **expected to score high** even on the most contested topics — that is the stance gate working as designed, not leniency.

Provenance & reproducibility

Every result carries provenance: the run ID, model, prompt version, a SHA-256 of the input text, and a UTC timestamp. The deterministic scorer ships with an 8-scenario self-test (`npm run selftest`), an offline demo (`npm run demo:bad`), and a live endpoint test (`npm run test:endpoint`).